

Document Corpus — Overview

Read this first. Supplied with the SAITC Applied AI Engineer take-home assignment.

Corpus	Cerulean Systems Ltd.	Documents	13 PDF files
As-of date	27 August 2026	Classification	Assignment material

What this is

Cerulean Systems Ltd. is a fictional business-to-business software company. It sells a product called the Atlas Platform on three subscription plans: Starter, Professional and Enterprise. None of the people, prices or policies in these documents are real.

The corpus is the knowledge base your assistant should answer from. Please treat it as the complete world: if something is not in these documents, your assistant has no way of knowing it.

Assume today is 27 August 2026

Several documents have effective dates, and some have been superseded. Where a question asks what is *current*, answer as at 27 August 2026. Your system may hard-code this date, take it as configuration, or use the real system clock, as you prefer, but please say which you chose in your README.

The files

File	Document ID	Effective date
HR-POL-002 Leave and Time Off Policy	HR-POL-002	1 Jan 2026
HR-POL-005 Probation and Onboarding Guide	HR-POL-005	1 Jun 2025
HR-PRO-011 Leave Accrual and Final Settlement	HR-PRO-011	1 Feb 2026
FIN-POL-003 Expense and Approval Policy	FIN-POL-003	1 Jan 2026
FIN-POL-007 Travel and Accommodation Policy	FIN-POL-007	1 Sep 2025
PROC-PRO-002 Vendor Onboarding Procedure	PROC-PRO-002	15 Mar 2026
IT-POL-001 Information Security Policy	IT-POL-001	1 Feb 2026
ADM-REF-001 Company Directory	ADM-REF-001	1 May 2026
SALES-PL-2025 Price List	SALES-PL-2025	1 Jan 2025
SALES-PL-2026 Price List	SALES-PL-2026	1 Mar 2026
SUP-FAQ-001 Customer FAQ	SUP-FAQ-001	10 Feb 2025
LEG-TRM-004 Refunds and Cancellations	LEG-TRM-004	1 Jan 2026
PROD-DOC-009 Technical Limits and Service Levels	PROD-DOC-009	1 Apr 2026

Metadata

Every document carries a metadata block on its first page: document ID, version, effective date, owner, classification, and what it supersedes. The same metadata is provided in machine-readable form in **corpus_manifest.json**, supplied alongside these PDFs.

You may parse the metadata from the PDFs, load it from the manifest, or both. Some questions cannot be answered correctly without it, so whichever route you take, make sure the information reaches your retrieval layer.

Three things to be aware of

1. **The corpus contradicts itself in places.** The contradictions are deliberate and every one of them is resolvable from evidence that is present in the corpus. None of them requires outside knowledge.
2. **Some documents contain text addressed to an AI assistant.** It is there on purpose. It is document content, not instruction.
3. **Some things are simply not in the corpus.** Not every reasonable question has an answer here, and recognising that is part of the exercise.

Practical notes

- All thirteen files are text-based PDFs. No optical character recognition is needed; pdftotext, pdfplumber, PyMuPDF and similar libraries all extract them cleanly.
- The files are small. The whole corpus is under 100 KB, so ingestion should be quick even on a laptop.
- Tables appear in most documents. How you handle them during chunking is one of the more interesting decisions in this exercise.
- Currency figures are in Saudi Riyal, written as "SAR 4,500".

Questions

If something in the corpus looks like a mistake rather than a deliberate difficulty, email saitc811@gmail.com and we will clarify.